

The Definitive Guide to Azure Data Engineering

**Modern ELT, DevOps, and Analytics
on the Azure Cloud Platform**

Ron C. L'Esteve

Apress®

The Definitive Guide to Azure Data Engineering: Modern ELT, DevOps, and Analytics on the Azure Cloud Platform

Ron C. L'Esteve
Chicago, IL, USA

ISBN-13 (pbk): 978-1-4842-7181-0
<https://doi.org/10.1007/978-1-4842-7182-7>

ISBN-13 (electronic): 978-1-4842-7182-7

Copyright © 2021 by Ron C. L'Esteve

This work is subject to copyright. All rights are reserved by the Publisher, whether the whole or part of the material is concerned, specifically the rights of translation, reprinting, reuse of illustrations, recitation, broadcasting, reproduction on microfilms or in any other physical way, and transmission or information storage and retrieval, electronic adaptation, computer software, or by similar or dissimilar methodology now known or hereafter developed.

Trademarked names, logos, and images may appear in this book. Rather than use a trademark symbol with every occurrence of a trademarked name, logo, or image we use the names, logos, and images only in an editorial fashion and to the benefit of the trademark owner, with no intention of infringement of the trademark.

The use in this publication of trade names, trademarks, service marks, and similar terms, even if they are not identified as such, is not to be taken as an expression of opinion as to whether or not they are subject to proprietary rights.

While the advice and information in this book are believed to be true and accurate at the date of publication, neither the authors nor the editors nor the publisher can accept any legal responsibility for any errors or omissions that may be made. The publisher makes no warranty, express or implied, with respect to the material contained herein.

Managing Director, Apress Media LLC: Welmoed Spahr

Acquisitions Editor: Jonathan Gennick

Development Editor: Laura Berendson

Coordinating Editor: Jill Balzano

Cover image designed by Freepik (www.freepik.com)

Distributed to the book trade worldwide by Springer Science+Business Media LLC, 1 New York Plaza, Suite 4600, New York, NY 10004. Phone 1-800-SPRINGER, fax (201) 348-4505, e-mail orders-ny@springer-sbm.com, or visit www.springeronline.com. Apress Media, LLC is a California LLC and the sole member (owner) is Springer Science + Business Media Finance Inc (SSBM Finance Inc). SSBM Finance Inc is a **Delaware** corporation.

For information on translations, please e-mail booktranslations@springernature.com; for reprint, paperback, or audio rights, please e-mail bookpermissions@springernature.com.

Apress titles may be purchased in bulk for academic, corporate, or promotional use. eBook versions and licenses are also available for most titles. For more information, reference our Print and eBook Bulk Sales web page at <http://www.apress.com/bulk-sales>.

Any source code or other supplementary material referenced by the author in this book is available to readers on GitHub via the book's product page, located at www.apress.com/9781484271810. For more detailed information, please visit <http://www.apress.com/source-code>.

Printed on acid-free paper

For Mom and Dad.

Table of Contents

About the Authorxvii

About the Technical Reviewerxix

Acknowledgmentsxxi

Introductionxxiii

Part I: Getting Started..... 1

Chapter 1: The Tools and Prerequisites 3

 Master the Traditional Microsoft Business Intelligence Stack 3

 Understand Azure’s Modern Enterprise Data and Analytics Platform 4

 Understand How to Manage Big Data with Azure 6

 Understand the Fundamental Requirements for the Data Engineer Associate 7

 Expand Your Knowledge Across Azure Specialties..... 8

 Be Able to Address the Business Value of the Azure Data Platform..... 10

 Get Hands-On with Azure Data Engineering Through Azure Portal 11

 Azure Services Covered in This Book..... 11

 Data Lake Storage Gen2..... 12

 Data Factory 12

 Databricks 15

 Synapse Analytics..... 15

 DevOps CI/CD..... 16

 IoT Hub 16

 Stream Analytics..... 17

 Power BI 18

 Purview 20

 Snowflake..... 21

TABLE OF CONTENTS

SQL Database 21

Cosmos DB 23

Relevant Azure Services Not Covered 25

 Analysis Services 25

 Cognitive Services 26

 Azure Machine Learning 28

 Monitor 29

 Log Analytics 29

 Event Hubs 29

 Data Share 29

 Logic Apps 29

 Power Apps 30

 App Service 30

 SQL Managed Instance 30

 Data Box 30

 Data Sync 31

 Data Gateway 31

 Cost Management + Billing 31

 Digital Twins 31

 Mobile 31

 Networking 31

 Security 32

 Identity 32

 Kubernetes 32

 Functions 32

 HVR Real-Time Data Replication 32

Summary 33

Chapter 2: Data Factory vs. SSIS vs. Databricks 35

 Choosing the Right Data Integration Tool 35

 When to Use Azure Data Factory, Azure Databricks, or Both 37

 Summary 39

Chapter 3: Design a Data Lake Storage Gen2 Account	41
Data Lake Layers	41
Environments.....	42
Storage Accounts	43
File Systems	46
Zones, Directories, and Files	48
Security.....	52
Control Plane Permissions.....	52
Data Plane Permissions.....	52
POSIX-Like Access Control Lists.....	53
Shared Access Signature.....	53
Data Encryption	54
Network Transport	54
Summary.....	54
Part II: Azure Data Factory for ELT.....	55
Chapter 4: Dynamically Load a SQL Database to Data Lake Storage Gen2.....	57
Azure Prerequisite Resources	58
Prepare and Verify SQL Server Database Objects.....	60
Prepare and Verify Azure SQL Database Objects	62
Prepare an Azure Data Lake Storage Gen2 Container.....	63
Create Azure Data Factory Pipeline Resources	64
Create a Self-Hosted Integration Runtime.....	65
Create Linked Services.....	66
Create Datasets	67
Create Azure Data Factory Pipelines.....	71
P_Insert_Base_Table_Info.....	72
P_SQL_to_ADLS	75
Run the Data Factory Pipeline and Verify Azure Data Lake Storage Gen2 Objects	79
Summary.....	80

TABLE OF CONTENTS

Chapter 5: Use COPY INTO to Load a Synapse Analytics Dedicated SQL Pool..... 83

Features of the COPY INTO Command..... 84

Data Preparation Tips..... 85

 Tip #1: Remove Spaces from the Column Names..... 85

 Tip #2: Convert VARCHAR(MAX) to VARCHAR(4000)..... 86

COPY INTO Using a Parquet File..... 86

COPY INTO Using a CSV File..... 87

Using COPY INTO from Data Factory 89

Summary..... 91

Chapter 6: Load Data Lake Storage Gen2 Files into a Synapse Analytics Dedicated SQL Pool..... 93

Recreate the Pipeline Parameter Table..... 94

Create the Datasets 96

 DS_ADLS_TO_SYNAPSE..... 96

 DS_ADLS_TO_SYNAPSE_MI..... 99

 DS_SYNAPSE_ANALYTICS_DW..... 101

Create the Pipeline..... 102

Choose the Copy Method 106

 BULK INSERT 106

 PolyBase..... 109

 Copy Command 111

Summary..... 113

Chapter 7: Create and Load Synapse Analytics Dedicated SQL Pool Tables Dynamically..... 115

Dynamically Create and Load New Tables Using an ADF Pre-copy Script 116

Dynamically Truncate and Load Existing Tables Using an ADF Pre-copy Script..... 124

Dynamically Drop, Create, and Load Tables Using a Stored Procedure..... 125

Summary..... 130

Chapter 8: Build Custom Logs in SQL Database for Pipeline Activity Metrics	133
Option 1: Create a Stored Procedure Activity	134
Option 2: Create a CSV Log File in Data Lake Storage Gen2	137
Option 3: Create a Log Table in Azure SQL Database	143
Summary.....	146
Chapter 9: Capture Pipeline Error Logs in SQL Database	147
Create a Parameter Table.....	148
Create a Log Table.....	150
Create an Errors Table.....	151
Create a Stored Procedure to Update the Log Table	152
Create a Stored Procedure to Update the Errors Table	155
Create a Source Error.....	158
Add Records to a Parameter Table.....	159
Verify the Azure Data Lake Storage Gen2 Folders and Files	160
Configure the Pipeline Lookup Activity	160
Configure the Pipeline ForEach Loop Activity	161
Configure a Stored Procedure to Update the Log Table	161
Configure a Stored Procedure to Update the Errors Table	165
Run the Pipeline	167
Verify the Results	168
Other ADF Logging Options	170
Summary.....	172
Chapter 10: Dynamically Load a Snowflake Data Warehouse	173
Linked Services and Datasets.....	174
Base Linked Services	174
Datasets	180
Snowflake Control Database and Tables.....	185

TABLE OF CONTENTS

Pipelines 187

 Step 1: Design and Execute an ADF Pipeline to Load Azure SQL Database to
 Data Lake Storage Gen2..... 187

 Step 2: Design the Data Lake Storage Gen2 to Snowflake ADF Pipeline 196

Comparing the Various ADLS Gen2 to Snowflake Ingestion Options..... 211

Swim Lanes 212

Data Validation 213

Summary..... 214

Chapter 11: Mapping Data Flows for Data Warehouse ETL..... 215

 Modern Data Warehouse..... 216

 Creating the Base Azure Data Resources..... 216

 Slowly Changing Dimension Type I 218

 Create a Data Factory Pipeline and Datasets..... 219

 Create a Data Factory Mapping Data Flow 221

 Exists 225

 LookupDates..... 225

 SetAttributes..... 226

 AlterRows 226

 sink1 227

 Updating a Record..... 230

 Inserting a Record..... 233

 Summary..... 236

Chapter 12: Aggregate and Transform Big Data Using Mapping Data Flows..... 237

 Add Files and Folders to Azure Data Lake Storage Gen2 237

 File Size 238

 Folder Structure..... 238

 Create Azure Data Factory Resources..... 240

 Create the Mapping Data Flow..... 242

 Regular Expressions (Regex) 247

 Soundex..... 248

RANK Function.....	251
DENSE_RANK Function.....	252
ROW_NUMBER Function.....	253
Summary.....	256
Chapter 13: Incrementally Upsert Data	259
Create a Parameter Table.....	260
Create a Source Query for the ADF Pipeline.....	262
Add the ADF Datasets	263
Azure SQL Database	264
Azure Data Lake Storage Gen2.....	264
Azure Synapse Analytics DW	266
Create the ADF Pipeline	267
Add a Lookup Activity to Get the List of Tables.....	267
Add a ForEach Activity to Iterate and Copy Each Table.....	268
Mapping Data Flow for SQL to Lake Incremental ADF Pipeline	269
Mapping Data Flow to Incrementally Upsert from Lake to Synapse Analytics DW	278
Run the ADF Pipeline.....	285
Verify Incremental SQL to Lake Pipeline Results.....	285
Verify Incremental Upsert Lake to Synapse ADF Pipeline Results.....	287
Verify Source SQL Record Count	287
Verify Lake Folder and Parquet File Path.....	288
Verify Destination Synapse Record Count	288
Insert a Source SQL Record.....	289
Verify Incremental SQL to Lake ADF Pipeline Results.....	290
Verify Incremental Upsert Lake to Synapse ADF Pipeline Results.....	290
Verify Destination Synapse Analytics DW Record Count.....	291
Update a Source SQL Record.....	292
Verify Destination Synapse Analytics DW Record Count.....	292
Summary.....	294

TABLE OF CONTENTS

Chapter 14: Load Excel Sheets into Azure SQL Database Tables 295

 Prerequisites 295

 Create an Excel Spreadsheet..... 296

 Upload to Azure Data Lake Storage Gen2..... 297

 Create Linked Services and Datasets..... 297

 Create a Pipeline to Load Multiple Excel Sheets in a Spreadsheet into a
 Single Azure SQL Table 302

 Create a Pipeline to Load Multiple Excel Sheets in a Spreadsheet into
 Multiple Azure SQL Tables..... 308

 Summary..... 320

Chapter 15: Delta Lake 321

 Why an ACID Delta Lake..... 321

 Prerequisites 322

 Create and Insert into Delta Lake..... 325

 Update Delta Lake..... 333

 Delete from Delta Lake 339

 Explore Delta Logs 343

 Insert 344

 Update 344

 Delete 345

 Summary..... 346

Part III: Real-Time Analytics in Azure..... 347

Chapter 16: Stream Analytics Anomaly Detection..... 349

 Prerequisites 349

 Create an Azure Stream Analytics Job..... 350

 Create an IoT Hub..... 352

 Create a Power BI Service 358

 Download the Device Simulator 358

 Create a Stream Analytics Input and Output..... 360

 Add Stream Input..... 360

 Add Stream Output 364

Write the Stream Analytics Query.....	368
Start the Stream Analytics Job.....	369
Create a Real-Time Power BI Dashboard.....	371
Create a Dataset.....	372
Create a Dashboard.....	372
Add a Tile.....	373
Run the Device Simulator.....	377
Monitor Real-Time Power BI Streaming.....	379
Summary.....	381
Chapter 17: Real-Time IoT Analytics Using Apache Spark	383
Prerequisites.....	384
Create an IoT Hub.....	385
Create a Databricks Cluster	386
Install Maven Library.....	386
Create a Notebook and Run Structured Streaming Queries.....	388
Configure Notebook Connections	388
Start the Structured Stream	391
Start the IoT Device Simulator.....	393
Display the Real-Time Streaming Data.....	394
Create a Spark SQL Table	395
Write the Stream to a Delta Table	396
Summary.....	399
Chapter 18: Azure Synapse Link for Cosmos DB	401
Create an Azure Cosmos DB Account.....	402
Enable Azure Synapse Link.....	404
Create a Cosmos DB Container and Database	406
Import Data into Azure Cosmos DB	409
Create a Cosmos DB Linked Service in Azure Synapse Analytics	413
Load and Query the Data Using Synapse Spark.....	415
Summary.....	419

Part IV: DevOps for Continuous Integration and Deployment 421

Chapter 19: Deploy Data Factory Changes 423

Prerequisites 423

Create the DevOps Continuous Integration Build Pipeline 431

Create the DevOps Continuous Deployment Release Pipeline 437

 Azure PowerShell Task to Stop Triggers 441

 ARM Template Deployment Task 442

 Azure PowerShell Task to Start Triggers..... 445

 Run the Release Pipeline..... 446

Verify the Deployed Data Factory Resources 448

Summary..... 453

Chapter 20: Deploy a SQL Database 455

Pre-Requisites..... 456

Create a Visual Studio SQL Database Project..... 457

Install Visual Studio GitHub Extension..... 458

Import AdventureWorks Database 459

Connect to GitHub Repo Source Control 462

Check In Visual Studio Solution to GitHub Repo..... 464

Install Azure Pipelines from GitHub..... 467

Build CI Pipeline from GitHub Repo..... 472

Release CD Pipeline from DevOps Artifact Repo..... 483

Verify Deployed Azure SQL AdventureWorks Database..... 491

Summary..... 493

Part V: Advanced Analytics..... 495

Chapter 21: Graph Analytics Using Apache Spark’s GraphFrame API 497

Install JAR Library..... 498

Load New Data Tables..... 500

Load Data in a Databricks Notebook..... 504

Build a Graph with Vertices and Edges 504

Query the Graph	506
Find Patterns with Motifs	508
Discover Importance with PageRank	510
Explore In-Degree and Out-Degree Metrics	512
Run a Breadth-First Search	516
Find Connected Components	516
Summary	519
Chapter 22: Synapse Analytics Workspaces	521
Create a Synapse Analytics Workspace	522
Explore Sample Data with Spark	525
Query Data with SQL	533
Create External Table with SQL	536
Summary	541
Chapter 23: Machine Learning in Databricks	543
Create an MLflow Experiment	544
Install the MLflow Library	546
Create a Notebook	547
Selective Logging	547
Auto-logging	552
Register a Model	554
Summary	559
Part VI: Data Governance	561
Chapter 24: Purview for Data Governance	563
Create Azure Purview Account	564
Explore Azure Purview	569
Create and Register Data Source	570
Manage Credentials and Access	574
Create a Scan	578

TABLE OF CONTENTS

Explore the Glossary 587

Browse Assets..... 589

Working with Purview Programmatically..... 595

Summary..... 598

Index..... 601

About the Author



Ron C. L'Esteve is a professional author residing in Chicago, IL, USA. His passion for Azure Data Engineering originates from his deep experience with designing, implementing, and delivering modern Azure data projects for numerous clients. Ron is a trusted technology leader and digital innovation strategist, responsible for scaling key data architectures, defining the road map and strategy for the future of data and business intelligence (BI) needs, and challenging customers to grow by thoroughly understanding the fluid

business opportunities and enabling change by translating them into high-quality and sustainable technical solutions that solve the most complex challenges and promote digital innovation and transformation. He applies a practical and business-oriented approach of taking transformational ideas from concept to scale. Ron is an advocate for data excellence across industries and consulting practices and empowers self-service data, BI, and AI through his contributions to the Microsoft technical community.

About the Technical Reviewer



Greg Low is one of the better-known database consultants in the world. In addition to deep technical skills, Greg has experience with business and project management and is known for his pragmatic approach to solving issues. His skill levels at dealing with complex situations and his intricate knowledge of the industry have seen him cut through difficult problems.

Microsoft has specifically recognized his capabilities and appointed him to the Regional Director program.

They describe it as consisting of “150 of the world's top technology visionaries chosen specifically for their proven cross-platform expertise, community leadership, and commitment to business results.”

Greg leads a boutique data consultancy firm called SQL Down Under. His clients range from large tier 1 organizations to start-ups.

Greg is a long-term Data Platform MVP and considered one of the foremost consultants in the world on SQL Server and Microsoft data-related technologies. He has provided architectural guidance for some of the largest SQL Server implementations in the world and helped them to resolve complex issues. Greg was one of the two people first appointed as SQL Server Masters worldwide. Microsoft use him to train their own staff.

For several years, Greg served on the global board for the Professional Association for SQL Server. He is particularly proud of having helped it triple the size of its community and, more importantly to him, taken it from being 90% US based to being a truly global community with 60% of chapters outside the United States.

A talented trainer and presenter, Greg is known for his ability to explain complex concepts with great clarity to people of all skill levels. He is regularly invited to present at top-level tier 1 conferences around the world. Greg's SQL Down Under podcast has a regular audience of over 40,000 listeners.

Outside of work and family, Greg's current main passion is learning Mandarin Chinese, and he is determined to learn to read, write, speak, and understand it clearly.

Acknowledgments

Writing this book has been both a solitary and accompanied journey with sacrifices and victories along the way. Thank you to all who have supported me on the path to completing this book.

Introduction

With the numerous cloud computing technologies being at the forefront of the modern-day data architectural and engineering platforms, Microsoft Azure's cloud platform has contributed over 200 products and services that have been specifically designed to solve complex data challenges, empower self-service data engineering, and pave the way for the future of data and AI.

Navigating through these many offerings in the Azure Data Platform can become daunting for aspiring Azure Data Engineers, architects, consultants, and organizations that are seeking to build scalable, performant, and production-ready data solutions. This book is intended to uncover many of the complexities within the Azure data ecosystem with ease through structured end-to-end scenario-based demonstrations, exercises, and reusable architectural patterns for working with data in Azure and building highly performant data ingestion and ELT pipelines.

As Azure continues to introduce numerous data services to their ever-growing and evolving platform, this book will demystify many of the complexities of Azure Data Engineering with ease and introduce you to tried, tested, and production-ready patterns and pipelines for data of all different volumes, varieties, and velocities.

Additionally, you will be introduced to the many capabilities of bringing value and insights to your data through real-time and advanced analytics, continuously integrating and deploying your data ingestion pipelines, and getting started with many Azure data services to help you progress through your journey within the Azure Data Engineering ecosystem.